

# Lab 4 - ML Programming

December 3, 2021

## 1 EXERCISE 0

### 1.1 Dataset Preprocessing

```
[1]: ## Import dataset and any libraries
## Convert any non-numeric values to numeric values

import numpy as np
import pandas as pd
import csv
import matplotlib.pyplot as plt

tictactoe = pd.read_csv('tic-tac-toe.data', header=None)
tictactoe.replace('x',0,inplace=True)
tictactoe.replace('o',1,inplace=True)
tictactoe.replace('b',2,inplace=True)
tictactoe.replace('negative',0,inplace=True)
tictactoe.replace('positive',1,inplace=True)
```

```
[2]: ## Confirm that the dataset is unbalanced

## Data imbalance refers to an unequal distribution of classes within a dataset
## So to show that the data is imbalanced we need the count of positive and
    ↪ negative endgames

tictactoe[9].value_counts()
```

```
[2]: 1    626
      0    332
      Name: 9, dtype: int64
```

```
[3]: ## Explain what stratified sampling and implement a stratified sampler

## Stratified sampling is a method of partitioning a population into
    ↪ subpopulations
## In this example we want a subpopulation which has an equal distribution
    ↪ across classes
```

```

## Pick the sampling number for each class
## Want to keep as much data as possible so make both classes around the size
↳ of the smaller class (i.e. 330)
n0 = 330
n1 = 330
## Get indexes and lengths for the classes respectively
idx0 = tictactoe.index.values[tictactoe[9] == 0]
idx1 = tictactoe.index.values[tictactoe[9] == 1]
len0 = len(idx0) ## 626
len1 = len(idx1) ## 332
## Draw randomly from the indices
draw0 = np.random.permutation(len0)[:n0]
idx0 = idx0[draw0]
draw1 = np.random.permutation(len1)[:n1]
idx1 = idx1[draw1]
## Combine the drawn indexes
idx = np.hstack([idx0, idx1])
## Create new dataset
new_tictactoe = tictactoe.loc[idx, :]
## Check length of new dataset (should be 2x332)
new_tictactoe[9].value_counts()

```

```

[3]: 0    330
     1    330
     Name: 9, dtype: int64

```

```

[4]: ## Split the data into a train(80%) and test(20%)

tictactoe_train = new_tictactoe.sample(frac=0.8, random_state=42) ## random
↳ state is just a seed value
tictactoe_test = new_tictactoe.drop(tictactoe_train.index)

```

## 2 EXERCISE 1

### 2.1 Logistic Regression w/ Gradient Descent

```

[5]: ## Split data into features and targets
x_train = tictactoe_train.iloc[:, :-1].values
y_train = tictactoe_train.iloc[:, -1].values
x_test = tictactoe_test.iloc[:, :-1].values
y_test = tictactoe_test.iloc[:, -1].values

```

```

[63]: ## Implement linear classification with stochastic gradient ascent algorithm
## Choose imax between 100 to 1000
## Use bolddriver as the step length controller
## In each iteration of the algorithm calculate |f(xi-1) - f(xi)| and logloss
↳ on test set

```

```

## Plot these against iteration number i and explain the graphs

## Code based on https://beckernick.github.io/logistic-regression-from-scratch/
def sigmoid(x):
    return 1.0/(1.0 + np.exp(-x))

def log_loss(y, p):
    return -(y * np.log(p) + (1 - y) * np.log(1 - p)).sum()

def logistic_regression(x, y, xTest, yTest, alpha, imax = 1000, precision = 0.
    ↪00000001):
    beta = np.zeros(x_train.shape[1])
    betaX = np.dot(beta, x.T)
    ## Calculating log likelihood
    log_likelihood = (y * betaX - np.log(1 + np.exp(betaX))).sum()
    ## For plotting
    iter_count = []
    absolute_dif = []
    loglosslist = []
    for i in range(0, imax):
        ## Calculate probability
        p = sigmoid(betaX)
        ## Add the gradient to maximize log likelihood
        gradient = np.dot(x.T, y - p)
        beta = beta + (alpha * gradient)
        betaX = np.dot(beta, x.T)
        log_likelihoodnew = (y * betaX - np.log(1 + np.exp(betaX))).sum()
        ## For plotting
        iter_count.append(i)
        absolute_dif.append(abs(log_likelihoodnew - log_likelihood))
        predictiontest = sigmoid(np.dot(beta, xTest.T))
        loglosslist.append(log_loss(yTest, predictiontest))
        ## Convergence condition
        if (abs(log_likelihoodnew - log_likelihood) < precision):
            return beta, iter_count, absolute_dif, loglosslist
        log_likelihood = log_likelihoodnew
        ## Bolddriver
        ## If the function values are decreasing, increase alpha. Otherwise, ↪
        ↪decrease alpha
        if log_likelihoodnew < log_likelihood:
            alpha = 1.05 * alpha
        else:
            alpha = 0.75 * alpha
    return beta, iter_count, absolute_dif, loglosslist

beta, xaxis, absolute_dif, loglosstest = logistic_regression(x_train, y_train, ↪
    ↪x_test, y_test, alpha=0.0001)

```

```


fig, axs = plt.subplots(1, 2, figsize=(16, 4))
axs[0].plot(xaxis, logloss_test, label = "iteration vs logLoss")
axs[0].grid()
axs[0].legend()

axs[1].plot(xaxis, absolute_dif, label = "iteration vs function difference")
axs[1].grid()
axs[1].legend()

## The logloss is increasing which is what we want as we are trying to maximize
→ our log
## The absolute difference as we update the betas is decreasing

```

[63]: <matplotlib.legend.Legend at 0x1b0813a05b0>



## 3 EXERCISE 2

### 3.1 Logistic Regression w/ Newton's Algorithm

```

[66]: ## Implement Newton Algorithm
## In each iteration calculate |f(xi-1)-f(xi)| and the logloss on test set
## Comment on the rate of convergence in the light of plots from above

def newton(x, y, xTest, yTest, alpha, imax = 1000, precision = 0.00000001):
    beta = np.zeros(x_train.shape[1])
    betaX = np.dot(beta, x.T)
    ## Calculating log likelihood
    log_likelihood = (y * betaX - np.log(1 + np.exp(betaX))).sum()
    ## For plotting
    iter_count = []
    absolute_dif = []
    loglosslist = []
    for i in range(0, imax):
        ## Calculate probability

```

```

p = sigmoid(betaX)
## Make Hessian matrix
w = np.diag(p*(1 - p))
hessian = x.T.dot(w).dot(x)
gradient = np.dot(x.T, y - p)
term = np.dot(np.linalg.inv(hessian),gradient)
beta = beta + (alpha * term)
betaX = np.dot(beta, x.T)
log_likelihoodnew = (y * betaX - np.log(1 + np.exp(betaX))).sum()
## For plotting
iter_count.append(i)
absolute_dif.append(abs(log_likelihoodnew - log_likelihood))
predictiontest = sigmoid(np.dot(beta, xTest.T))
loglosslist.append(log_loss(yTest, predictiontest))
## Convergence condition
if (abs(log_likelihoodnew - log_likelihood) < precision):
    return beta, iter_count, absolute_dif, loglosslist
log_likelihood = log_likelihoodnew
return beta, iter_count, absolute_dif, loglosslist

beta, xaxis, absolute_dif, loglosstest = newton(x_train, y_train, x_test,
→y_test, alpha=0.0001)



fig, axs = plt.subplots(1, 2, figsize=(16, 4))
axs[0].plot(xaxis, loglosstest, label = "iteration vs logLoss")
axs[0].grid()
axs[0].legend()

axs[1].plot(xaxis, absolute_dif, label = "iteration vs function difference")
axs[1].grid()
axs[1].legend()

## Since figure 3 tells us to use a fixed step length, we get straight lines
→instead of curves
## However, just like the graphs using gradient ascent, we see that the
→absolute difference decreases
## This means that the difference between the functions decreases as our betas
→are optimized
## We can see that it is much slower in convergence here
## However, if we add in a steplength controller like bolddriver, convergence
→is quicker with Newton
## See following cell for graphs

```

[66]: <matplotlib.legend.Legend at 0x1b083841430>



```
[57]: ## If we redo task 2 implementing the bold driver like in task 1, we get very
      → similar graphs
      ## However, the Newton algorithm seems to converge faster than gradient ascent

def newton_bold(x, y, xTest, yTest, alpha, imax = 1000, precision = 0.00000001):
    beta = np.zeros(x_train.shape[1])
    betaX = np.dot(beta, x.T)
    ## Calculating log likelihood
    log_likelihood = (y * betaX - np.log(1 + np.exp(betaX))).sum()
    ## For plotting
    iter_count = []
    absolute_dif = []
    loglosslist = []
    for i in range(0, imax):
        ## Calculate probability
        p = sigmoid(betaX)
        ## Make Hessian matrix
        w = np.diag(p*(1 - p))
        hessian = x.T.dot(w).dot(x)
        gradient = np.dot(x.T, y - p)
        term = np.dot(np.linalg.inv(hessian), gradient)
        beta = beta + (alpha * term)
        betaX = np.dot(beta, x.T)
        log_likelihoodnew = (y * betaX - np.log(1 + np.exp(betaX))).sum()
        ## For plotting
        iter_count.append(i)
        absolute_dif.append(abs(log_likelihoodnew - log_likelihood))
        predictiontest = sigmoid(np.dot(beta, xTest.T))
        loglosslist.append(log_loss(yTest, predictiontest))
        ## Convergence condition
        if (abs(log_likelihoodnew - log_likelihood) < precision):
            return beta, iter_count, absolute_dif, loglosslist
        log_likelihood = log_likelihoodnew
    ## Bold driver
```

```

    ## If the function values are decreasing, increase alpha. Otherwise,
    ↪ decrease alpha
    if log_likelihoodnew < log_likelihood:
        alpha = 1.05 * alpha
    else:
        alpha = 0.75 * alpha
    return beta, iter_count, absolute_dif, loglosslist

beta, xaxis, absolute_dif, loglosstest = newton_bold(x_train, y_train, x_test, ↪
    ↪ y_test, alpha=0.0001)

fig, axs = plt.subplots(1, 2, figsize=(16, 4))
axs[0].plot(xaxis, loglosstest, label = "iteration vs logLoss")
axs[0].grid()
axs[0].legend()

axs[1].plot(xaxis, absolute_dif, label = "iteration vs function difference")
axs[1].grid()
axs[1].legend()

```

[57]: <matplotlib.legend.Legend at 0x1b085203100>

