

Lab 5 - ML Programming

December 10, 2021

1 EXERCISE 1

1.1 Backward search for variable selection

```
[85]: ## Import data and necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

bank = pd.read_csv('bank.csv', delimiter=";")

[86]: ## Convert any non-numeric values to numeric values
## Remove NA/missing values

## Recode month and y columns as numbers to reduce sparsity
bank['month'] = bank['month'].
    ↳ replace(['jan', 'feb', 'mar', 'apr', 'may', 'jun', 'jul', 'aug', 'sep', 'oct', 'nov', 'dec'],
            ↳ ['1', '2', '3', '4', '5', '6', '7', '8', '9', '10', '11', '12'])
bank['y'] = bank['y'].replace(['no', 'yes'], ['0', '1'])

## Separate numeric and categorical columns
numerical_columns = ['age',
    ↳ 'balance', 'day', 'month', 'duration', 'campaign', 'pdays', 'previous', 'y']
cat_cols = [c for c in bank.columns if c not in numerical_columns]

## Converting columns that are numerical into numeric dtypes
for col in numerical_columns:
    bank[col] = bank[col].astype(np.float32)

bank_numeric = pd.concat([bank[numerical_columns],
    ↳ pd.get_dummies(bank[cat_cols])], axis=1)
bank_numeric = bank_numeric.dropna()

[77]: bank_numeric['y'].value_counts()

[77]: 0.0    4000
      1.0    521
```

Name: y, dtype: int64

```
[78]: ## Balance the dataset, keeping as much data as possible
n0 = 520
n1 = 520
## Get indexes and lengths for the classes respectively
idx0 = bank_numeric.index.values[bank_numeric['y'] == 0.0]
idx1 = bank_numeric.index.values[bank_numeric['y'] == 1.0]
len0 = len(idx0) ## 4000
len1 = len(idx1) ## 521
## Draw randomly from the indices
draw0 = np.random.permutation(len0)[:n0]
idx0 = idx0[draw0]
draw1 = np.random.permutation(len1)[:n1]
idx1 = idx1[draw1]
## Combine the drawn indexes
idx = np.hstack([idx0, idx1])
## Create new dataset
bank_numeric = bank_numeric.loc[idx, :]
## Check length of new dataset (should be 2x520)
bank_numeric['y'].value_counts()
```

```
[78]: 0.0    520
      1.0    520
Name: y, dtype: int64
```

```
[79]: ## Split the data into a train/test splits according to the ratios 80%:20%
bank_numeric_train = bank_numeric.sample(frac=0.8, random_state=42) ## random
      ↪ state is just a seed value
bank_numeric_test = bank_numeric.drop(bank_numeric_train.index)
```

```
[80]: ## Normalize the data with  $xi - \mu$ 
## We should first normalize the training data
for column in bank_numeric_train.columns:
    bank_numeric_train[column] = ↪
    ↪ (bank_numeric_train[column] - bank_numeric_train[column].mean()
    ↪ ) / bank_numeric_train[column].std()

## To normalize test set, apply normalization parameters obtained from training
↪ set
for column in bank_numeric_test.columns:
    bank_numeric_test[column] = ↪
    ↪ (bank_numeric_test[column] - bank_numeric_train[column].mean()
    ↪ ) / bank_numeric_train[column].std()
```

```
[ ]: ## Implement logistic regression and mini-batch Gradient Ascent
## Keep learning rate and batch size fixed and iteratively do backward
→selection keeping track of the AIC metric
## Report the final error on test set

## Split data into features and targets
x_train = bank_numeric_train.iloc[:, :-1].values
y_train = bank_numeric_train.iloc[:, -1].values
x_test = bank_numeric_test.iloc[:, :-1].values
y_test = bank_numeric_test.iloc[:, -1].values

def sigmoid(x, beta):
    return 1.0 / (1.0 + np.exp(-1 * (np.dot(x, beta))))

def gradient(x, y, beta):
    return np.dot(x.T, (y - sigmoid(x, beta)))

## Code based on https://stackoverflow.com/questions/38157972/
→how-to-implement-mini-batch-gradient-descent-in-python
def create_batch(x, y, batch_size, shuffle=False):
    assert x.shape[0] == y.shape[0]
    if shuffle:
        indices = np.arange(x.shape[0])
        np.random.shuffle(indices)
    for start_idx in range(0, x.shape[0] - batch_size + 1, batch_size):
        if shuffle:
            excerpt = indices[start_idx: start_idx + batch_size]
        else:
            excerpt = slice(start_idx, start_idx + batch_size)
        yield x[excerpt], y[excerpt]

def mini_batch_GA(x, y, batch_size=42, alpha = 0.01, imax = 1000, precision = 0.
    →00000001):
    beta = np.zeros((x.shape[1], 1))
    for i in range(0, imax):
        mini_batches = create_batch(x, y, batch_size)
        for mini_batch in mini_batches:
            x_mini, y_mini = mini_batch
            beta = beta + alpha * gradient(x_mini, y_mini, beta)
    return beta

def log_likelihood(x, y, beta):
    return (y * betaX - np.log(1 + np.exp(betaX))).sum()

def AIC(x, y, beta, p):
    return 2 * p - 2 * log_likelihood(x, y, beta)
```

```

def log_loss(y, p):
    return -(y * np.log(p) + (1 - y) * np.log(1 - p)).mean()

def backward_search(x,y, xTest, yTest):
    v_used = [[i] for i in range(x.shape[1])]
    improvement = True
    iterations = 0
    AIClist = []
    loglosslist = []
    while(improvement):
        iterations += 1
        AIClist.append(AIC(x,y,mini_batch_GA(x,y),len(v_used)))
        gain_best = 0
        v_best = []
        for v in v_used:
            ## Calculate AIC for x without v
            x_without_v = np.delete(x, v, 1)
            beta_without_v = mini_batch_GA(x_without_v,y)
            AIC_without_v = AIC(x_without_v,y,beta_without_v, len(v_used))
            ## Calculate AIC for x with v
            x_with_v = np.insert(x, v, 1, axis=1)
            beta_with_v = mini_batch_GA(x_with_v,y)
            AIC_with_v = AIC(x_with_v,y,beta_with_v, len(v_used))
            gain = AIC_without_v - AIC_with_v
            v_used_v = [i[0] for i in v_used if i != v]
            if gain > gain_best:
                gain_best = gain
                v_best = v
        improvement = True if gain_best > 0 else False
        if improvement:
            v_used = [i for i in v_used if i != v_best]
            AIClist.append(gain_best)
            predictiontest = sigmoid(np.dot(beta, xTest.T))
            loglosslist.append(log_loss(yTest, predictiontest))
            finalerror = loglosslist[-1]
    return iterations, AIClist, finalerror

iterations, AIClist, finalerror = backward_search(x_train,y_train,x_test,y_test)
print(finalerror)

plt.plot(iterations, AIClist, label = "iteration vs AIC")
plt.grid()
plt.legend()

```

2 EXERCISE 2

2.1 Regularization for Logistic Regression

```
[ ]: ## Pick a range of  $\alpha$  and  $\lambda$  defined on grid
## You can choose fixed batchsize = 50.
## Implement k-fold cross-validation protocol for grid search
## For each combination of  $\alpha$  and  $\lambda$  you will perform k-fold cross-validation.
    → Let  $k = 5$  in this case
## Keep track of mean performance (i.e. Classification Accuracy value) across  $k$ 
    → folds for each set of hyperparameters
## Plot on the grid  $\alpha$  vs  $\lambda$  the Classification Accuracy score for all
    → combinations
## For the optimal value of  $\alpha$  and  $\lambda$ , train your model on complete training
    → data and evaluate on Test data
## Report one single Accuracy and Log-likelihood for Test data
## Plot Train and Validation Accuracy and Log-likelihood metrics per  $k$  - fold
    → iteration

imax = 100
alpha_list = np.array([0.01, 0.1, 1])
Lambda_list = np.array([0.1, 1, 10])
batch_size = 50
k = 5

def logreg_kfold(x, y, xTest, yTest, imax, k, alpha_list, Lambda_list,
    → batch_size):
    beta = np.zeros(x.shape[1])
    ## Y prediction with lambda
    y_hat = lambda x, beta : (np.exp(x@beta)/(np.exp(x@beta)+1)).reshape(-1,1)
    def sigmoid(x,beta):
        return 1.0/(1.0 + np.exp(-1*(np.dot(x,beta))))
    def gradient(x, y, beta):
        return np.dot(x.T,(y-sigmoid(x, beta)))
    def log_likelihood(x,y,beta):
        return (y * betaX - np.log(1 + np.exp(betaX))).sum()
    def create_batch(x,y,batch_size, shuffle=False):
        assert x.shape[0] == y.shape[0]
        if shuffle:
            indices = np.arange(x.shape[0])
            np.random.shuffle(indices)
        for start_idx in range(0, x.shape[0] - batch_size + 1, batch_size):
            if shuffle:
                excerpt = indices[start_idx:start_idx + batch_size]
            else:
                excerpt = slice(start_idx, start_idx + batch_size)
        yield x[excerpt], y[excerpt]
```

```

def mini_batch_GA(x, y, batch_size=42, alpha = 0.01, imax = 1000, precision_
↪= 0.00000001):
    beta = np.zeros((x.shape[1], 1))
    for i in range(0,imax):
        mini_batches = create_batch(x, y, batch_size)
        for mini_batch in mini_batches:
            x_mini, y_mini = mini_batch
            beta = beta + alpha * gradient(x_mini, y_mini, beta)
    return beta

## Grid
Grid = np.zeros((len(alpha_list)*len(Lambda_list),k))
## K-folds setting
index = np.random.permutation(x.index)
for i in range(1, k+1):
    x_val = x.loc[index][((i-1)* int(np.floor(x.shape[0]/k)):i* int(np.
↪floor(x.shape[0]/k)))]
    x_train = x.drop(index = x_val.index)
    y_val = y.loc[x_val.index]
    y_train = y.drop(index = x_val.index)
    Grid_counter = 0
    ## Alpha Loop
    for alpha_counter in range(len(alpha_list)):
        alpha = alpha_list[alpha_counter]
        ## Lambda Loop
        for Lambda_counter in range(len(Lambda_list)):
            Lambda = Lambda_list[Lambda_counter]
            ## Error vectors setting
            error_train = np.zeros(imax)
            error_val = np.zeros(imax)
            error_test = np.zeros(imax)
            ## Iterations
            for j in range(1,imax+1):
                ## Logistic regression setup
                index = np.random.permutation(x_train.index)
                ## Iterations of log reg
                for u in range(1, int(np.ceil(len(x_train)/batch_size))+1):
                    x_train_i = x_train.loc[Index][((i-1)*batch_size:
↪i*batch_size)]
                    y_train_i = y_train.loc[Index][((i-1)*batch_size:
↪i*batch_size)]

                    beta = mini_batch_GA(x_train_i,y_train_i)
                    ## Log likelihood computation
                    error_train[j-1] = log_likelihood(x_train,y_train,beta)
                    error_val[j-1] = log_likelihood(x_val,y_val,beta)
                Grid[Grid_counter,i-1] = error_val[imax-1]
            Grid_counter += 1

```

```
logreg_kfold(x_train, x_test, y_train, y_test, imax, k, alpha_list,
↳Lambda_list, batch_size)
```

3 EXERCISE 3

3.1 Implementing hyperband

```
[82]: ## Firstly, re-split the dataset into train/validation/test splits according to
↳ratios: 70% : 15% : 15%
```

```
bank_num_train = bank_numeric.sample(frac=0.7,random_state=42) ## random state
↳is just a seed value
bank_num_leftover = bank_numeric.drop(bank_num_train.index)
bank_num_validation = bank_num_leftover.sample(frac=0.5,random_state=42)
bank_num_test = bank_num_leftover.drop(bank_num_validation.index)

print(bank_num_train.shape,bank_num_validation.shape,bank_num_test.shape)
```

```
(728, 41) (156, 41) (156, 41)
```

```
[ ]: ## Implement hyperband algorithm
```

```
## Code taken from https://homes.cs.washington.edu/~jamieson/hyperband.html
```

```
max_iter = 100 ## maximum iterations/epochs per configuration
eta = 3 ## defines downsampling rate (default=3)
logeta = lambda x: log(x)/log(eta)
s_max = int(logeta(max_iter)) ## number of unique executions of Successive
↳Halving (minus one)
B = (s_max+1)*max_iter ## total number of iterations (without reuse) per
↳execution of Successive Halving (n,r)
```

```
## Begin Finite Horizon Hyperband outterloop. Repeat indefinitely.
```

```
for s in reversed(range(s_max+1)):
    n = int(ceil(int(B/max_iter/(s+1))*eta**s)) ## initial number of
↳configurations
    r = max_iter*eta**(-s) ## initial number of iterations to run
↳configurations for

    ## Begin Finite Horizon Successive Halving with (n,r)
    T = [ get_random_hyperparameter_configuration() for i in range(n) ]
    for i in range(s+1):
        ## Run each of the n_i configs for r_i iterations and keep best n_i/eta
        n_i = n*eta**(-i)
        r_i = r*eta**(i)
        val_losses = [
↳run_then_return_val_loss(num_iters=r_i,hyperparameters=t) for t in T ]
```

```
T = [ T[i] for i in argsort(val_losses)[0:int( n_i/eta )] ]  
## End Finite Horizon Successive Halving with (n,r)
```