


# Lab 6 - ML Programming

December 17, 2021

## 1 EXERCISE 0

### 1.1 Data pre-processing

```
[67]: ## Generate a Sample dataset called D1 :  
## Initialize matrix x R100×1 using Uniform distribution with  $\mu = 1$  and  $\sigma = 0.05$   
## Generate target y R100×1 using  $y = 1.3x^2 + 4.8x + 8 + \epsilon$ , where  $\epsilon$  R100×1  
## Wine Quality called D2: (use winequality-red.csv)  
  
import matplotlib.pyplot as plt  
import numpy as np  
import pandas as pd  
from sklearn import datasets, linear_model  
from sklearn.metrics import mean_squared_error, r2_score  
from sklearn.model_selection import GridSearchCV  
from sklearn.model_selection import cross_val_score  
from sklearn.preprocessing import PolynomialFeatures  
from sklearn.linear_model import LinearRegression  
from sklearn.linear_model import Ridge  
  
D1_x = np.random.normal(loc=1, scale=0.05, size=(100,1))  
random = np.random.rand(100,1)  
D1_y = (1.3*(D1_x**2)) + (4.8*D1_x) + 8 + random  
  
D2 = pd.read_csv('winequality-red.csv',delimiter=';')  
  
plt.scatter(D1_x,D1_y)  
plt.show()
```



## 2 EXERCISE 1

### 2.1 Generalized Linear Models with Scikit Learn

```
[12]: ## Split wine data into test and train using 80%-20% split
D2_train = D2.sample(frac=0.8,random_state=42) ## random state is just a seed
    ↪ value
D2_test = D2.drop(D2_train.index)
```

```
[13]: ## Normalize the data with  $x_i - \mu /$ 
## We should first normalize the training data
for column in D2_train.columns:
    D2_train[column] = (D2_train[column]-D2_train[column].mean())/
    ↪ D2_train[column].std()

## To normalize test set, apply normalization parameters obtained from training
    ↪ set
for column in D2_test.columns:
    D2_test[column] = (D2_test[column]-D2_train[column].mean())/
    ↪ D2_train[column].std()
```

```
[14]: ## Split data into features and targets
x_train = D2_train.iloc[:, :-1].values
y_train = D2_train.iloc[:, -1].values
x_test = D2_test.iloc[:, :-1].values
```

```
y_test = D2_test.iloc[:, -1].values
```

```
[16]: ## ORDINARY LEAST SQUARES
## Pick 3 sets of hyperparameters and learn each model (without cross
validation)
## Measure Train and Test RMSE and plot it on one plot
## Explain the plots and relate it to the theory studied in lectures


OLS_alphas = [0.01,0.05,0.1]
RMSE_train = []
RMSE_test = []

def calc_rmse(y, predictions):
    mse = mean_squared_error(y, predictions)
    rmse = np.sqrt(mse)
    return rmse

for i in OLS_alphas:
    model = linear_model.SGDRegressor(loss='squared_loss', penalty='none',
    alpha=0, \
                                shuffle=True, learning_rate='constant', eta0=i)
    model.fit(x_train,y_train)
    y_pred = model.predict(x_train)
    y_testpred = model.predict(x_test)
    RMSE_train.append(calc_rmse(y_train,y_pred))
    RMSE_test.append(calc_rmse(y_test,y_testpred))

plt.plot(OLS_alphas, RMSE_train, 'green', label='Training RMSE')
plt.plot(OLS_alphas, RMSE_test, 'blue', label='Test RMSE')
plt.title('OLS Training and Test RMSE')
plt.xlabel('OLS_alphas')
plt.ylabel('RMSE')
plt.legend()
plt.grid()
plt.show()

## Plot shows that as the learning rate increases, the RMSE increases as well
## The dramatic increase on test but not on training RMSE suggests the model is
overfitting for high learning rates
```



```
[18]: ## RIDGE REGRESSION
## Pick 3 sets of hyperparameters and learn each model (without cross_
validation)
## Measure Train and Test RMSE and plot it on one plot
## Explain the plots and relate it to the theory studied in lectures


ridge_alphas = [[0.001,0.001],[0.01,0.01],[0.05,0.05]]
ridge_RMSE_train = []
ridge_RMSE_test = []

for rows,cols in ridge_alphas:
    ridge_model = linear_model.SGDRegressor(loss='squared_loss', penalty='l2',
    alpha=cols, \
                                shuffle=True, learning_rate='constant',
    eta0=rows)
    ridge_model.fit(x_train,y_train)
    y_pred_ridge = ridge_model.predict(x_train)
    y_testpred_ridge = ridge_model.predict(x_test)
    ridge_RMSE_train.append(calc_rmse(y_train,y_pred_ridge))
    ridge_RMSE_test.append(calc_rmse(y_test,y_testpred_ridge))

plt.plot(ridge_RMSE_train, label = 'Training RMSE')
plt.plot(ridge_RMSE_test, label = 'Test RMSE')
```

```
plt.title('Ridge Training and Test RMSE')
plt.legend()
plt.grid()
plt.show()
```

*## Here we can see that the gap between the training and test RMSE is much larger*  
*## This suggests that the regularization term in ridge regression contributes more to overfitting*



```
[19]: ## LASSO
## Pick 3 sets of hyperparameters and learn each model (without cross-validation)
## Measure Train and Test RMSE and plot it on one plot
## Explain the plots and relate it to the theory studied in lectures

lasso_alphas = [[0.001,0.001],[0.01,0.01],[0.05,0.05]]
lasso_RMSE_train = []
lasso_RMSE_test = []

for rows,cols in lasso_alphas:
    lasso_model = linear_model.SGDRegressor(loss='squared_loss', penalty='l1',
alpha=cols, \
```

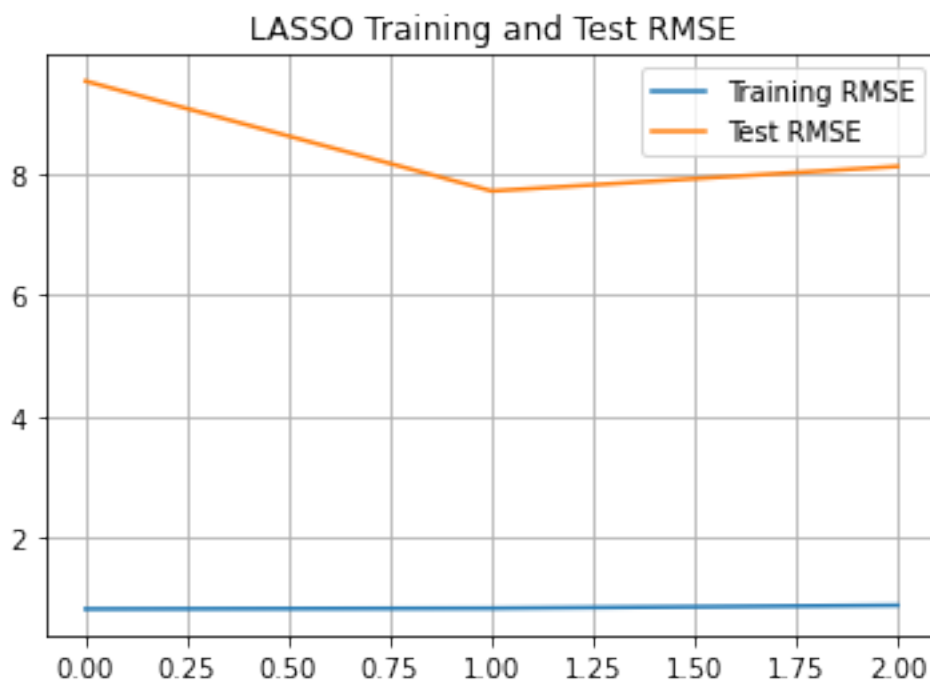
```

shuffle=True, learning_rate='constant',
eta0=rows)
lasso_model.fit(x_train,y_train)
y_pred_lasso = lasso_model.predict(x_train)
y_testpred_lasso = lasso_model.predict(x_test)
lasso_RMSE_train.append(calc_rmse(y_train,y_pred_lasso))
lasso_RMSE_test.append(calc_rmse(y_test,y_testpred_lasso))

plt.plot(lasso_RMSE_train, label = 'Training RMSE')
plt.plot(lasso_RMSE_test, label = 'Test RMSE')
plt.title('LASSO Training and Test RMSE')
plt.legend()
plt.grid()
plt.show()

## Plot shows us something similar to the ridge regression
## Test RMSE is much higher due to the penalty fitting to the training data

```



```

[20]: ## ORDINARY LEAST SQUARES
## Tune the hyperparameters using scikit learn GridSearchCV
## Plot the results of cross validation

hp = {'eta0':[0.01,0.05,0.1]}


```

```

model = linear_model.SGDRegressor(loss='squared_loss', penalty='none', alpha=0,
    ↪\
                                shuffle=True, learning_rate='constant')
model_gridsearch = GridSearchCV(model, hp, cv=5, return_train_score=True)
model_gridsearch.fit(x_train,y_train)

plt.plot(hp['eta0'], model_gridsearch.cv_results_['mean_train_score'], 'red',
    ↪label='Training Set')
plt.plot(hp['eta0'], model_gridsearch.cv_results_['mean_test_score'],'blue',
    ↪label='Test Set')
plt.xlabel('Learning Rate')
plt.ylabel('mean score')
plt.legend()
plt.grid()
plt.show()

```



```

[21]: ## ORDINARY LEAST SQUARES
      ## Cross validation

model = linear_model.SGDRegressor(loss='squared_loss', penalty='none', alpha=0,
    ↪\
                                shuffle=True, learning_rate='constant',
    ↪eta0=model_gridsearch.best_params_['eta0'])

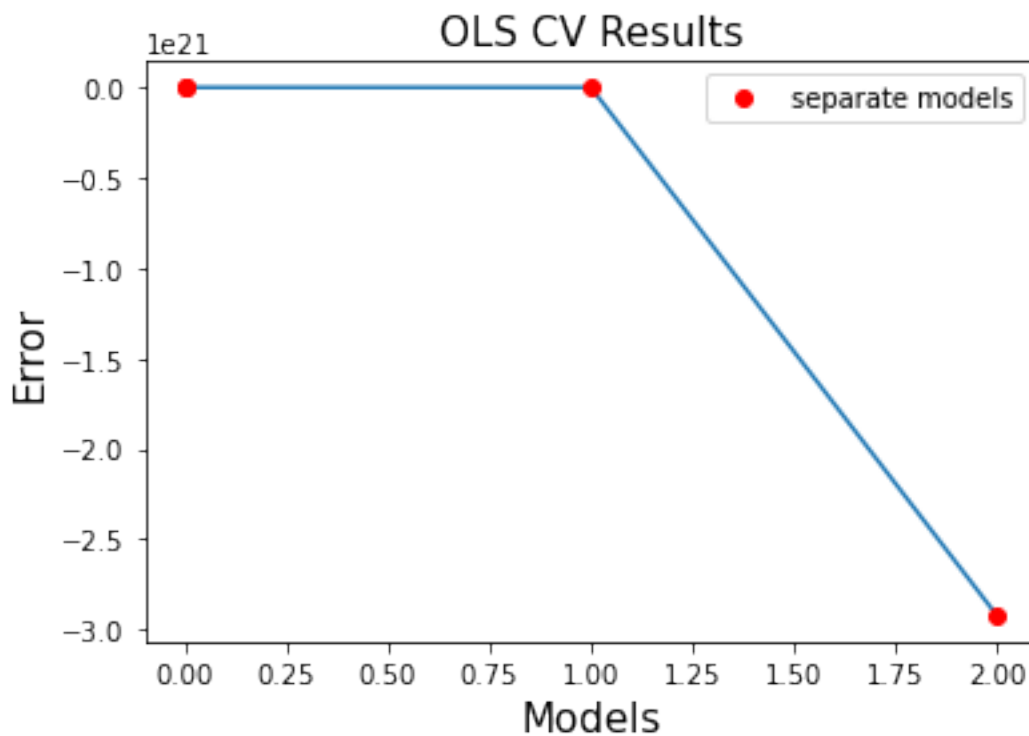
```

```

model.fit(x_train,y_train)
cv_train = cross_val_score(model,x_train,y_train,cv=4)
cv_test = cross_val_score(model,x_test,y_test,cv=4)

plt.title("OLS CV Results", fontsize = 15)
plt.plot(model_gridsearch.cv_results_["mean_test_score"])
plt.plot(model_gridsearch.cv_results_["mean_test_score"], "ro", label = "↪separate models")
plt.plot(model_gridsearch.best_score_, "ro")
plt.xlabel('Models', fontsize = 15)
plt.ylabel('Error', fontsize = 15)
plt.legend()
plt.show()

```



```

[22]: ## RIDGE REGRESSION
      ## Tune the hyperparameters using scikit learn GridSearchCV
      ## Plot the results of cross validation for each model

      hp = {'eta0':[0.01,0.05,0.1], 'alpha':[0.01,0.05,0.1]}
      ridge = linear_model.SGDRegressor(loss='squared_loss', penalty='l2', \
                                         shuffle=True, learning_rate='constant')
      ridge_gridsearch = GridSearchCV(ridge, hp, cv=5, return_train_score=True)

```

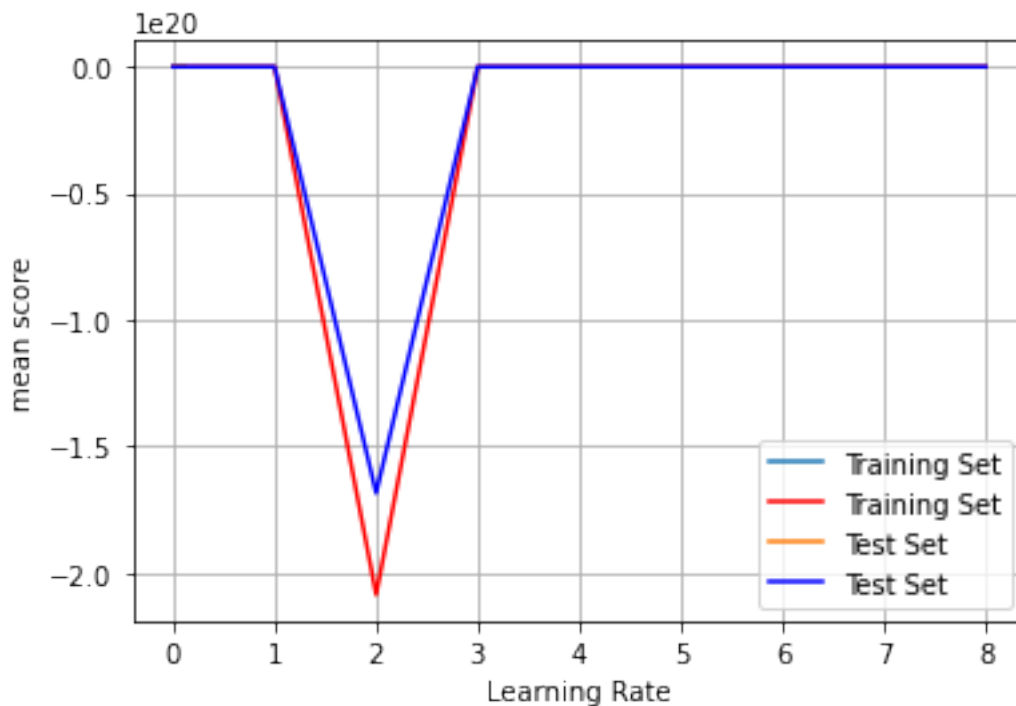


```

ridge_gridsearch.fit(x_train,y_train)

plt.plot(hp['eta0'], hp['alpha'], ridge_gridsearch.
    ↪cv_results_['mean_train_score'], 'red', label='Training Set')
plt.plot( hp['eta0'], hp['alpha'], ridge_gridsearch.
    ↪cv_results_['mean_test_score'],'blue', label='Test Set')
plt.xlabel('Learning Rate')
plt.ylabel('mean score')
plt.legend()
plt.grid()
plt.show()

```



```

[23]: ## RIDGE REGRESSION
      ## Cross validation


ridge = linear_model.SGDRegressor(loss='squared_loss', penalty='l2',
    ↪alpha=ridge_gridsearch.best_params_['alpha'], \
                                shuffle=True, learning_rate='constant',
    ↪eta0=ridge_gridsearch.best_params_['eta0'])
ridge.fit(x_train,y_train)
ridge_cv_train = cross_val_score(ridge,x_train,y_train,cv=4)
ridge_cv_test = cross_val_score(ridge,x_test,y_test,cv=4)

```

```

plt.title("Ridge Regression CV Results", fontsize = 15)
plt.plot(ridge_gridsearch.cv_results_["mean_test_score"])
plt.plot(ridge_gridsearch.cv_results_["mean_test_score"], "ro", label = "separate models")
plt.plot(ridge_gridsearch.best_score_, "ro")
plt.xlabel('Models', fontsize = 15)
plt.ylabel('Error', fontsize = 15)
plt.legend()
plt.show()

```



```

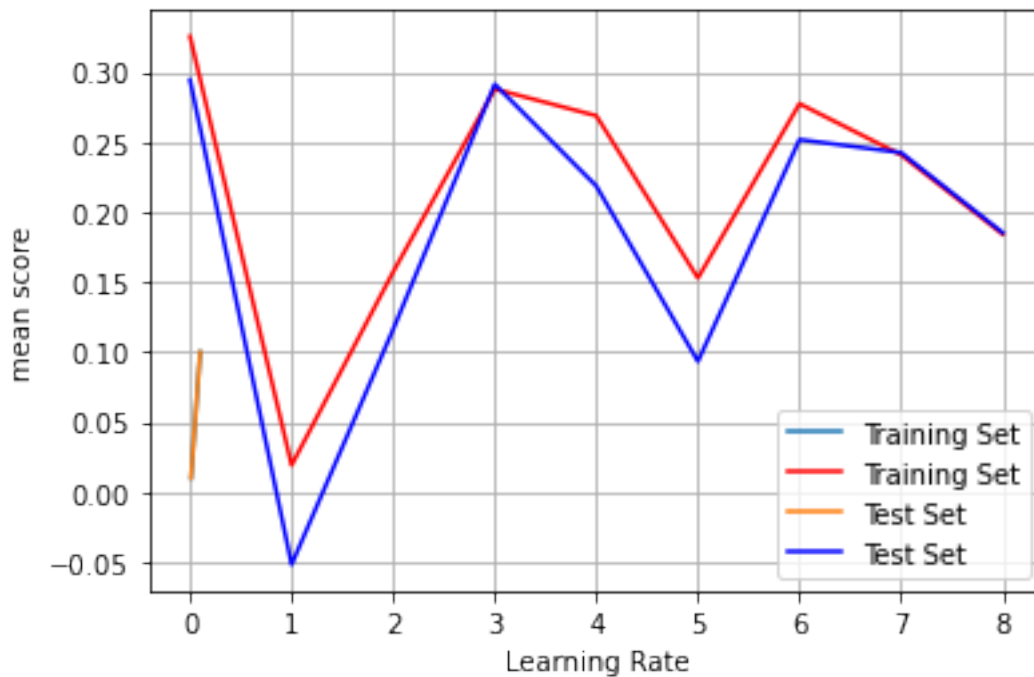
[27]: ## LASSO
      ## Tune the hyperparameters using scikit learn GridSearchCV
      ## Plot the results of cross validation for each model

      hp = {'eta0':[0.01,0.05,0.1], 'alpha':[0.01,0.05,0.1]}
      lasso = linear_model.SGDRegressor(loss='squared_loss', penalty='l1', \
                                         shuffle=True, learning_rate='constant')
      lasso_gridsearch = GridSearchCV(lasso, hp, cv=5, return_train_score=True)
      lasso_gridsearch.fit(x_train,y_train)

      plt.plot(hp['eta0'], hp['alpha'], lasso_gridsearch.
               ↪cv_results_['mean_train_score'], 'red', label='Training Set')

```

```
plt.plot( hp['eta0'], hp['alpha'], lasso_gridsearch.
    ↪cv_results_['mean_test_score'],'blue', label='Test Set')
plt.xlabel('Learning Rate')
plt.ylabel('mean score')
plt.legend()
plt.grid()
plt.show()
```

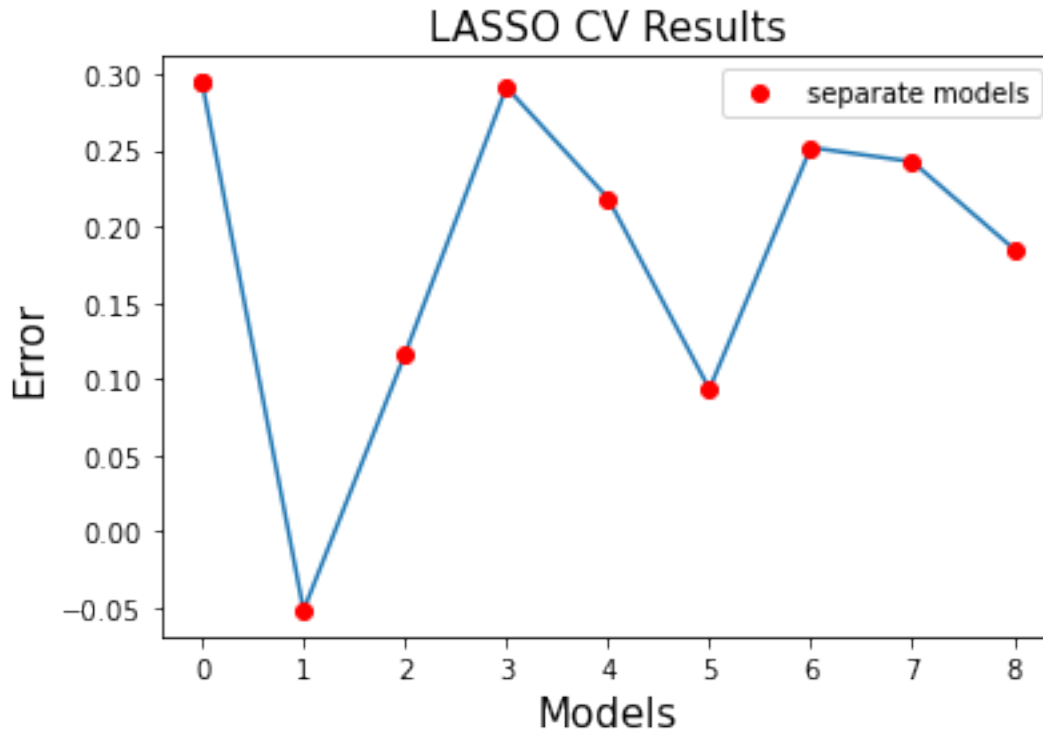


```
[28]: ## LASSO
      ## Cross validation

lasso = linear_model.SGDRegressor(loss='squared_loss', penalty='l1',
    ↪alpha=lasso_gridsearch.best_params_['alpha'], \
                                   shuffle=True, learning_rate='constant',
    ↪eta0=lasso_gridsearch.best_params_['eta0'])
lasso.fit(x_train,y_train)
lasso_cv_train = cross_val_score(lasso,x_train,y_train,cv=4)
lasso_cv_test = cross_val_score(lasso,x_test,y_test,cv=4)

plt.title("LASSO CV Results", fontsize = 15)
plt.plot(lasso_gridsearch.cv_results_["mean_test_score"])
plt.plot(lasso_gridsearch.cv_results_["mean_test_score"], "ro", label =
    ↪"separate models")
plt.plot(lasso_gridsearch.best_score_, "ro")
```

```
plt.xlabel('Models', fontsize = 15)
plt.ylabel('Error', fontsize = 15)
plt.legend()
plt.show()
```



```
[30]: ## Using the optimal hyperparameter you have to evaluate each model on the Test Set
      ##Report the results in a meaningful manner

      print(f"OLS best score is {model_gridsearch.best_score_}")
      print(f"Ridge Regression best score is {ridge_gridsearch.best_score_}")
      print(f"LASSO best score is {lasso_gridsearch.best_score_}")

      ## Based on these we see that LASSO is the best model based on scores generated by grid search cv
```

```
OLS best score is 0.29090210832334124
Ridge Regression best score is 0.2840577158597672
LASSO best score is 0.294460661890987
```

## 3 EXERCISE 2

### 3.1 Higher Order Polynomial Regression

```
[68]: ## Use higher degrees of polynomial feature for your data i.e. degrees 1, 2, 7, 10, 16 and 100
      ## TASK A
      ## For each newly created dataset learn LinearRegression
      ## Plot the predicted curves for each dataset


      ## Plot for original data points
      fig, axs = plt.subplots(1,1,figsize=(10,5))
      axs.scatter(D1_x, D1_y, color='red', label="data")
      axs.legend()

      degrees = [1, 2, 7, 10, 16,100]

      ## Plot for models with different polynomial features
      for i in degrees:
          poly = PolynomialFeatures(degree=i)
          poly_x = poly.fit_transform(D1_x)
          linreg = LinearRegression().fit(poly_x, D1_y)
          y_pred = linreg.predict(poly_x)
          concat = zip(*sorted(zip(D1_x,y_pred)))
          x,y = concat
          axs.plot(x, y, label=f"degree {i}", linewidth=3)
      axs.legend()
      plt.suptitle('Prediction with high degree of polynomials',fontsize=15)
      plt.show()

      ## In this plot we see that the higher the polynomial features, the more the
      line fits to the data points
      ## This is why models with high degrees of freedom tend to overfit the data
      ## Lines with too little degrees tend to underfit the data due to
      oversimplification of the relationship
      ## Therefore we need to find a good balance between the two
```

Prediction with high degree of polynomials



```
[69]: ## TASK B
## Fix the degree of polynomial to 10
## Pick 4 values of (regularization constant) and learn Ridge Regression
## Plot the predicted curves for each dataset
## Explain the phenomena you observed for different prediction curves


## Plot for original data points
fig, axs = plt.subplots(1,1,figsize=(10,5))
axs.scatter(D1_x, D1_y, color='red', label="data")
axs.legend()

## Lambda values given in problem
lambdas = [0, (10**-6), (10**-2), 1]

for i in lambdas:
    poly = PolynomialFeatures(degree=10)
    poly_x = poly.fit_transform(D1_x)
    ridge = Ridge(alpha=i).fit(poly_x, D1_y)
    y_pred = ridge.predict(poly_x)
    concat = zip(*sorted(zip(D1_x,y_pred)))
    x,y = concat
    axs.plot(x, y, label=f"lambda = {i}", linewidth=3)
axs.legend()
plt.suptitle('Effect of regularization',fontsize=15)
plt.show()
```

```
## Using regularization we see that there is much less likelihood of
→ overfitting in this case
## When lambda = 0 this represents no regularization and we see that it is more
→ prone to being skewed by noise in data
```

Effect of regularization



## 4 EXERCISE 3

### 4.1 Implementing Coordinate Descent

```
[ ]: ## Implement the Coordinate Descent algorithm
## Maintain a history of your values
## After training plot them against iterations in a single plot

## Reference: https://xavierbourretsicotte.github.io/lasso_implementation.
→html#Implementing-coordinate-descent-for-lasso-regression-in-Python
def coordinate_descent(x,y,itters):
    m,n = x.shape
    theta = np.zeros((n,1))
    theta_list = []
    for i in range(itter):
        for j in range(len(theta)):
            xtheta = np.dot(x,theta)
            update = ((y-xtheta).T).dot(x[:,j]) / (x[:,j].T).dot(x[:,j])
            theta_list.append(update)
        theta[i+1] = theta[j]
    return theta, theta_list
```

```

[77]: ##Implement CD with L1 regularization
## Maintain a history of your values
## After training plot them against iterations in a single plot

## Soft threshold constitutes part of the closed-form LASSO solution using L1
→regularization
def soft_threshold(rho,lamda):
    if rho < - lamda:
        return (rho + lamda)
    elif rho > lamda:
        return (rho - lamda)
    else:
        return 0

def coordinate_descent_lasso(x,y,lamda = .01, num_iters=100):
    m, n = x.shape
    theta = np.ones((n,1))
    theta_list = []
    for i in range(num_iters):
        for j in range(n):
            x_j = x[:,j].reshape(-1,1)
            y_pred = x @ theta
            rho = x_j.T @ (y - y_pred + theta[j]*x_j)
            theta[j] = soft_threshold(rho, lamda)
    return theta.flatten()

m,n = x_train.shape
initial_theta = np.ones((n,1))
theta_list = list()

theta = coordinate_descent_lasso(x_train,y_train, num_iters=100)
theta_list.append(theta)

#Stack into numpy array
theta_lasso = np.stack(theta_list).T

#Plot results
n,_ = theta_lasso.shape
plt.figure(figsize = (12,8))

for i in range(n):
    plt.plot(lamda, theta_lasso[i])

plt.xscale('log')

```



```
plt.xlabel('Log($\\lambda$)')
plt.ylabel('Coefficients')
plt.title('Lasso Paths - Numpy implementation')
plt.legend()
plt.axis('tight')
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-77-f0dc1e433d80> in <module>
    28 theta_list = list()
    29
---> 30 theta = coordinate_descent_lasso(x_train,y_train, num_iters=100)
    31 theta_list.append(theta)
    32

<ipython-input-77-f0dc1e433d80> in coordinate_descent_lasso(x, y, lamda,
-> num_iters)
    21         y_pred = x @ theta
    22         rho = x_j.T @ (y - y_pred + theta[j]*x_j)
---> 23         theta[j] = soft_threshold(rho, lamda)
    24     return theta.flatten()
    25

ValueError: could not broadcast input array from shape (1279,) into shape (1,)
```

```
[ ]: ## Compare the plots of the unregularized and regularized CD
      ## Highlight the difference
      ## What information can be inferred from these values?
```