

Lab 10 - ML Programming

January 29, 2022

1 EXERCISE 1

1.1 Exploring Movie Recommendation Dataset

```
[55]: ## Import data + necessary libraries

import matplotlib.pyplot as plt
import matplotlib as mp
import numpy as np
import pandas as pd
import math
plt.style.use('ggplot')

data = pd.read_csv('u.data',delimiter="\t", names=['user id','movie_
    ↳id','rating', 'timestamp'])
user = pd.read_csv('u.user',delimiter="|",names=['user_
    ↳id','age','gender','occupation','zip code'])
item = pd.read_csv('u.item',delimiter="|",names=['movie id','movie_
    ↳title','release date','video release date',
        'IMDb URL','unknown','Action','Adventure','Animation',
        'Children's','Comedy','Crime','Documentary','Drama','Fantasy',
        'Film-Noir','Horror','Musical','Mystery','Romance','Sci-Fi',
        'Thriller','War','Western'])
item.drop(['video release date','IMDb URL'], axis=1, inplace= True)
## Just lists of genre and occupation, not as relevant
genre = pd.read_csv('u.genre',delimiter="|",header=None)
genre.drop(genre.columns[1], axis=1, inplace=True)
genre.columns = ['Genres']
genre_list = list(genre['Genres'])
occupation = pd.read_csv('u.occupation',delimiter="|",header=None)
```

```
[2]: ## Showcase how the ratings vary across users
## Consider whether the plot is able to tell if most ratings are only from a_
    ↳handful of users

## Merge datasets into one dataframe
```

```

data_user = pd.merge(data[['user id','movie id','rating','timestamp']],
    ↳user[['user id','age','gender','occupation']], on='user id')
full = pd.merge(data_user[['user id','movie id',
    ↳id','rating','timestamp','age','gender','occupation']], item[['movie id',
    ↳id','movie title','release date',
        'unknown','Action','Adventure','Animation',
        "Children's",'Comedy','Crime','Documentary','Drama','Fantasy',
        'Film-Noir','Horror','Musical','Mystery','Romance','Sci-Fi',
        'Thriller','War','Western']], on='movie id')


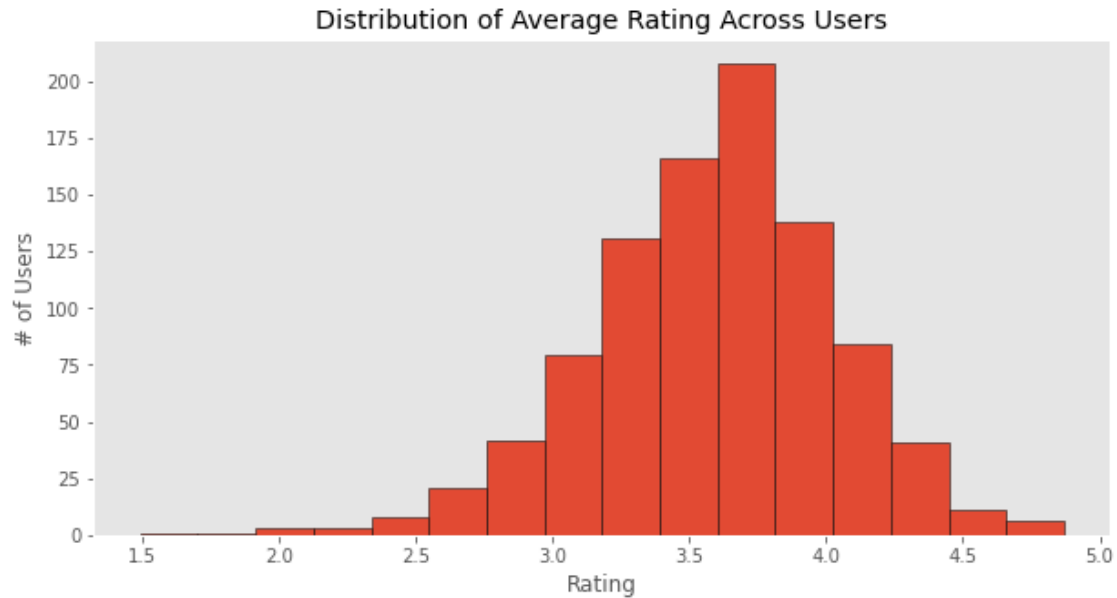
## Merged the data so that we get each users average rating and it's not just a
    ↳handful of users
## Shows us that most users have similar average ratings
## Color graph in a gradient in order to differentiate between high and low
    ↳averages
dftmp = full[['user id','rating']].groupby('user id').mean()
users = list(range(1,944))
dftmp['user id'] = users
dftmp['rating_dup'] = dftmp['rating']

dftmp2 = full[['user id','rating']].groupby('user id').mean()
dftmp2.hist(bins=16, grid=False, edgecolor='black', figsize=(10,5))
plt.xlabel('Rating')
plt.ylabel('# of Users')
plt.title('Distribution of Average Rating Across Users')
plt.show

plt.figure(figsize=(15, 5))
plt.scatter(dftmp['user id'],y=dftmp['rating'],c=dftmp['rating_dup'],cmap=plt.
    ↳cm.plasma)
plt.xlabel('User ID')
plt.ylabel('Average Rating')
plt.title('Average Rating Across Users')
plt.show

```

[2]: <function matplotlib.pyplot.show(close=None, block=None)>




```
[3]: ## Showcase how the ratings vary across items

dftmp = full[['movie id', 'rating']].groupby('movie id').mean()

dftmp.hist(bins=16, grid=False, density=True, edgecolor='black', figsize=(10,5))
for gen in genre_list:
    dftmp = full[full[gen]==1]
    dftmp = dftmp[['movie id', 'rating']].groupby('movie id').mean()
    dftmp['rating'].plot(grid=False, kind='kde', alpha=0.6, label=gen)
plt.legend(loc=(1.05,0), ncol=2)
plt.xlim(0,5)
plt.xlabel('Rating')
```

```
plt.ylabel('Movie Density')
plt.title("Distribution of ratings across movies")
plt.show()
```

We see that movies overall are rated around 3
 ## Film-noir genre is often rated at 3.5 and unknown doesn't have enough data
 ↳ to generate a normal bell curve
 ## However we can see that the distribution across genres is quite similar
 ↳ besides these 2 categories



[43]: ## Are there genres that are more highly rated than others?

```
dftmp_list = []
for gen in genre_list:
    dftmp = full[full[gen] == 1]
    dftmp = dftmp[['movie id', 'rating']].groupby('movie id').mean()
    dftmp[gen] = dftmp['rating']
    dftmp.drop(['rating'], axis=1, inplace=True)
    dftmp_list.append(dftmp)

dftmp = pd.concat(dftmp_list, ignore_index=True)
dicts = {}
for gen in genre_list:
    dicts[gen] = dftmp[gen].mean()
genre_avgrating = pd.DataFrame.from_dict(dicts, orient='index')
genre_avgrating['genre'] = genre_list
genre_avgrating.sort_values(by=[0], inplace=True, ascending=False)
plt.figure(figsize=(15, 5))
plt.bar(genre_avgrating['genre'], genre_avgrating[0], edgecolor='black')
plt.xlabel('Genres')
plt.xticks(rotation=70)
```

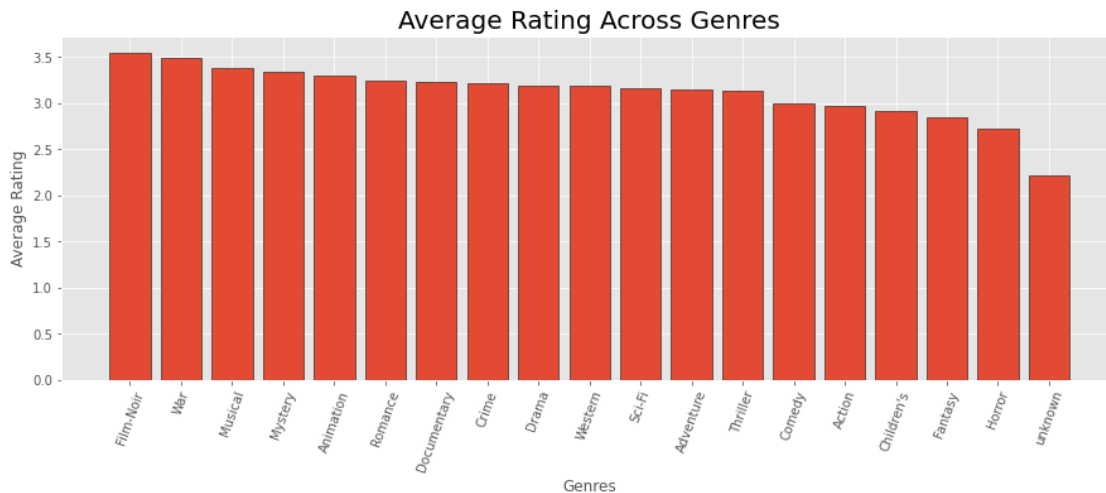
```

plt.ylabel('Average Rating')
plt.title('Average Rating Across Genres',fontsize=20)
plt.show


dictz = {}
for gen in genre_list:
    dictz[gen] = full[gen].mean()
proportion = pd.DataFrame.from_dict(dictz,orient='index')
proportion['genre'] = genre_list
proportion.sort_values(by=[0],inplace=True,ascending=True)
plt.figure(figsize=(10, 10))
plt.pie(proportion[0], labels=proportion['genre'], startangle=0,
    ↳rotatelabels=True)
plt.title('Proportion of Movies in Each Genre',fontsize=20)
plt.show()

## We see that film noir is the highest rated but it makes up an extremely
↳small proportion of the movies overall
## The more movies in a genre have been rated, the more reliable the average
↳rating
## This is why it is important to analyze these two graphs in conjunction

```



Proportion of Movies in Each Genre



```
[54]: ## What age groups prefer what genres based on ratings?
      ## You can bin respective ages to your preference

      bins= [0,10,18,30,50,200]
      labels = ['<=10','<=18','<=30','<=50','50+']
      full['age group'] = pd.cut(full['age'], bins=bins, labels=labels, right=True)


      for l in labels:
          age_grouped = full[full['age group'] == l]
          dftmp_list = []
          for gen in genre_list:
              dftmp = age_grouped[age_grouped[gen] == 1]
              dftmp = dftmp[['movie id','rating']].groupby('movie id').mean()
```

```


    dftmp[gen] = dftmp['rating']
    dftmp.drop(['rating'], axis=1, inplace= True)
    dftmp_list.append(dftmp)
dftmp = pd.concat(dftmp_list,ignore_index=True)
dicts = {}
for gen in genre_list:
    dicts[gen] = dftmp[gen].mean()
genre_avgrating = pd.DataFrame.from_dict(dicts,orient='index')
genre_avgrating['genre'] = genre_list
genre_avgrating.sort_values(by=[0],inplace=True,ascending=False)
plt.figure(figsize=(10, 5))
plt.bar(genre_avgrating['genre'],genre_avgrating[0],edgecolor='black')
plt.xlabel('Genres')
plt.xticks(rotation=70)
plt.ylabel('Average Rating')
plt.title(f'Average Rating Across Genres for Ages {l}',fontsize=20)
plt.show
dictz = {}
for gen in genre_list:
    dictz[gen] = age_grouped[gen].mean()
proportion = pd.DataFrame.from_dict(dictz,orient='index')
proportion['genre'] = genre_list
proportion.sort_values(by=[0],inplace=True,ascending=True)
plt.figure(figsize=(6, 6))
plt.pie(proportion[0], labels=proportion['genre'], startangle=0,
↪rotatelabels=True)
plt.title(f'Proportion of Movies in Each Genre for Ages f{l}',fontsize=20)
plt.show()

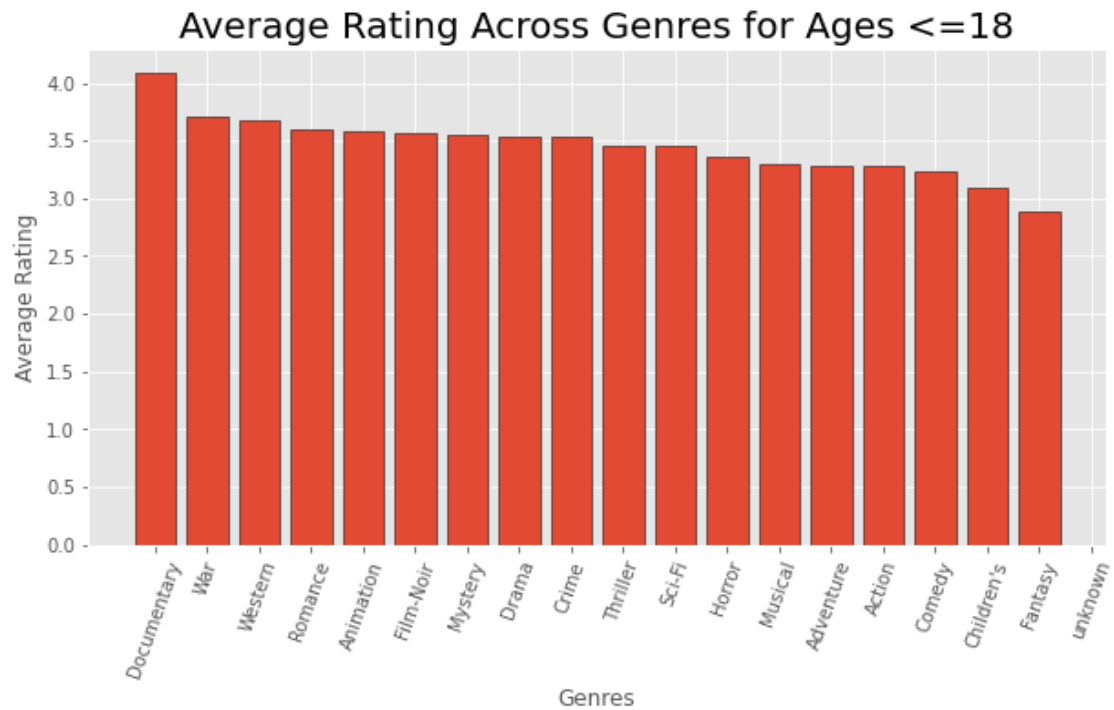
## This shows us the highest average rated genres for each age group first
## Then shows us the most frequently watched genres for each age group
## Should try to combine these measurements into one for most useful insights
## (i.e. normalize for amount of ratings per genre)

```





Proportion of Movies in Each Genre for Ages ≤ 10







Proportion of Movies in Each Genre for Ages ≤ 18



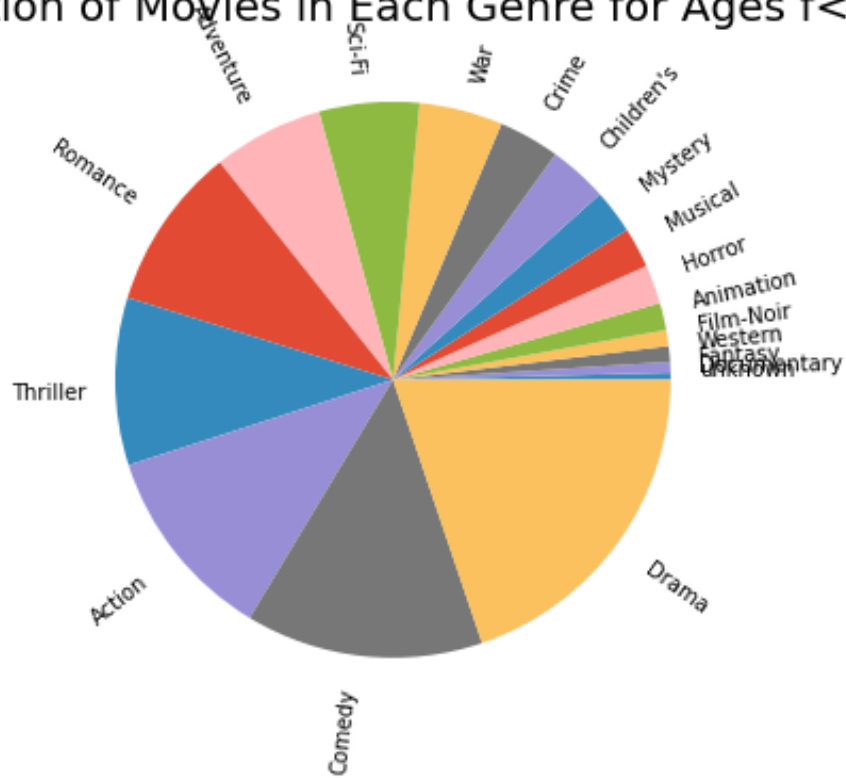


Proportion of Movies in Each Genre for Ages $f \leq 30$






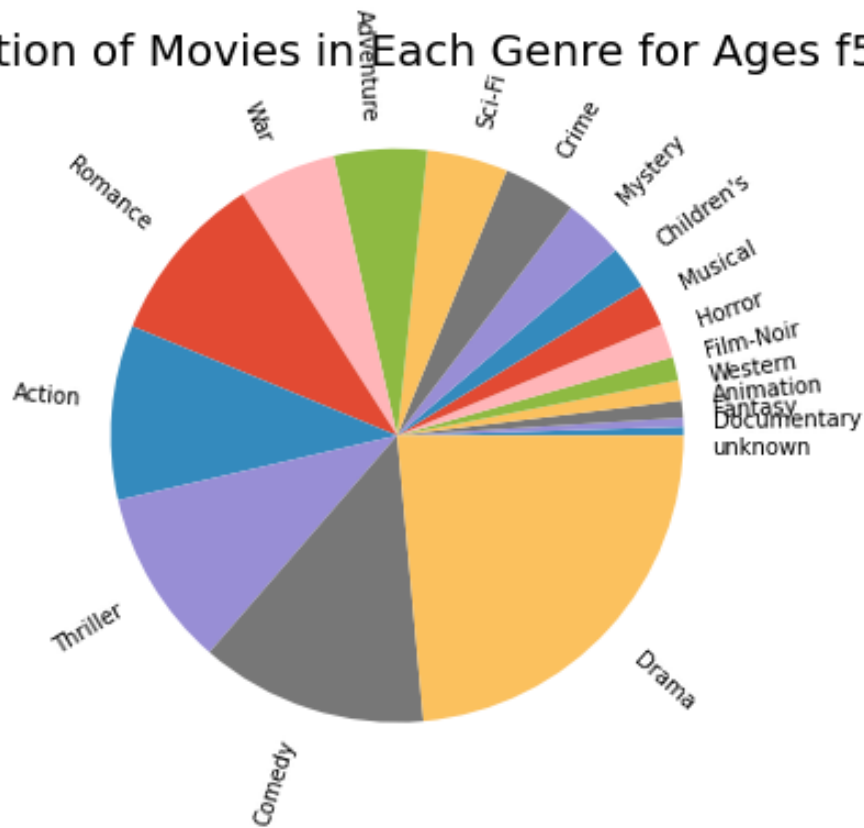
Proportion of Movies in Each Genre for Ages $f \leq 50$



Average Rating Across Genres for Ages 50+



Proportion of Movies in Each Genre for Ages f50+



2 EXERCISE 2

2.1 Implementing basic matrix factorization (MF) technique for recommender systems

```
[78]: ## Split data into 80-10-10

train = full.sample(frac=0.8,random_state=3116)
ytrain = train['rating']
xtrain = train.drop(['rating'], axis=1, inplace= True)
leftover = full.drop(train.index)
validation = leftover.sample(frac=0.5,random_state=3116)
yval = validation['rating']
xval = validation.drop(['rating'],axis=1, inplace= True)
test = leftover.drop(validation.index)
ytest = test['rating']
xtest = test.drop(['rating'],axis=1, inplace= True)
```

```
print(train.shape, validation.shape, test.shape)
```

```
(80000, 28) (10000, 28) (10000, 28)
```

```
[74]: def RMSE(y, y_hat):
    rmse = ((np.sum((y-y_hat)**2))/(len(y)))*0.5
    return rmse

## Based on https://github.com/Quang-Vinh/matrix-factorization
def obtain_rating(P, Q, users, items, ubias, ibias, bias):
    pred = bias + ubias.loc[users] + ibias.loc[item] + P.loc[users].dot(Q.loc[items]).
    ↪T)
    return pred

def sgd(sample, P, Q, ubias, ibias, bias, alpha, beta):
    for i, j, r in sample:
        pred = obtain_rating(P, Q, i, j, ubias, ibias, bias)
        error = r - pred
        ubias.loc[i] += alpha*(e[0] - beta*ubias.loc[i])
        ibias.loc[j] += alpha*(e[0] - beta*ibias.loc[i])
        P.loc[i] += alpha*(e[0]*Q.loc[j] - beta*P.loc[i])
        Q.loc[j] += alpha*(e[0]*P.loc[i] - beta*Q.loc[j])
    return P, Q, ubias, ibias

def mf(matrix, K, alpha, beta, epoch, sample, tsample):
    users, items = matrix.shape
    P = pd.DataFrame(np.random.normal(scale=1./K,
    ↪size=(users, k)), columns=[0, 1], index=sample['user id'].unique())
    Q = pd.DataFrame(np.random.normal(scale=1./K,
    ↪size=(items, k)), columns=[0, 1], index=sample['movie id'].unique())
    ubias = pd.DataFrame(0, columns=[0], index=sample['user id'].unique())
    ibias = pd.DataFrame(0, columns=[0], index=sample['movie id'].unique())
    sample = sample.values
    matrix = matrix.values
    bias = np.mean(matrix[np.where(matrix!=0)])
    rmse = []
    for i in range(epoch):
        np.random.shuffle(sample)
        P, Q, ubias, ibias = sgd(sample, P, Q, ubias, ibias, bias, alpha, beta)
        err = RMSE(y, y_hat)
        rmse.append(err)
    return rmse, P, Q, ubias, ibias, bias
```

```
[82]: def predict(matrix, test):
    y_hat = []
    values = []
    for i in range(0, len(test)):

```

```

marker = 1
if(test.iloc[i]['user id'] in matrix.index):
    if(test.iloc[i]['movie id'] in matrix.columns):
        marker = 0
    if(marker):
        values.append(i)
    else:
        y_hat.append(matrix[test.iloc[i]['movie id']][test.iloc[i]['user_
↪id']])
    return y_hat,test.drop(test.index[values])

rmse,P,Q,ubias,ibias,bias = mf()

```

```

-----
TypeError                                Traceback (most recent call last)
<ipython-input-82-297f64c16c54> in <module>
    20     return y_hat,test.drop(test.index[values])
    21
----> 22 nP, nQ = matrix_factorization(train_array, P, Q, k_latent)
    23 y_hat,new = predict(validation,xval)
    24 rmse_validation = rmse(yval,y_hat)

<ipython-input-74-cbfad11f3244> in matrix_factorization(R, P, Q, K, steps,
↪alpha, beta)
     8     for i in range(len(R)):
     9         for j in range(len(R[i])):
----> 10             if R[i][j] > 0:
    11                 eij = R[i][j] - np.dot(P[i,:],Q[:,j])
    12                 for k in range(K):

TypeError: '>' not supported between instances of 'str' and 'int'

```

3 EXERCISE 3

3.1 Recommender Systems using matrix factorization sklearn

```

[ ]: def divide(data,k):
    size_of_k = math.floor(len(data)/k)
    data_kdivided = []
    c = 0
    for i in range(0,k):
        dataset = pd.DataFrame(data.head(0))
        for j in range(i*size_of_k,(i*size_of_k)+size_of_k):
            dataset = dataset.append(data.iloc[j])
            c = c+1
        data_kdivided.append(dataset)

```

```

    for j in range(c,len(data)):
        data_kdivided[k-1]=data_kdivided[k-1].append(data.iloc[j])
    return data_kdivided

def k_traintest(x,k):
    kfold_data = []
    for i in range(0,k):
        xtest = x[i]
        xtrain = pd.DataFrame()
        for j in range(0,k):
            if i != j:
                xtrain = xtrain.append(x[j])
        final = dict([('xtrain',xtrain),('xtest',xtest)])
        kfold_data.append(final)
    return kfold_data

def kfold(xtrain,k,alpha,lamda):
    errors = []
    xtrain_k = divide(xtrain,k)
    kdata = k_traintest(xtrain_k,k)
    for i in range(0,k):
        matrixx = kdata[i]['xtrain'].pivot(index='user id',columns='movie_
→id',values='rating')
        matrixx = matrixx.replace(np.nan,0)
        rmse,P,Q,ubias,ibias,bias =_
→mf(matrixx,2,alpha,lamda,1,kdata[i]['xtrain'][['user id','movie_
→id','rating']],kdata[i]['xtest'][['user id','movie id','rating']])
        errors.append(sum(rmse)/len(rmse))
    return sum(errors)/k

```

```

[72]: rmse = kfold(train,3,0.01,0.01)
      rmse

```

```

-----
TypeError                                Traceback (most recent call last)
~\anaconda3\lib\site-packages\pandas\core\ops\array_ops.py in_
→_na_arithmetic_op(left, right, op, is_cmp)
    141     try:
--> 142         result = expressions.evaluate(op, left, right)
    143     except TypeError:

~\anaconda3\lib\site-packages\pandas\core\computation\expressions.py in_
→evaluate(op, a, b, use_numexpr)
    234         # error: "None" not callable
--> 235         return _evaluate(op, op_str, a, b) # type: ignore[misc]
    236     return _evaluate_standard(op, op_str, a, b)

```



```

~\anaconda3\lib\site-packages\pandas\core\computation\expressions.py in
↳ _evaluate_numexpr(op, op_str, a, b)
    119     if result is None:
--> 120         result = _evaluate_standard(op, op_str, a, b)
    121

```

```

~\anaconda3\lib\site-packages\pandas\core\computation\expressions.py in
↳ _evaluate_standard(op, op_str, a, b)
    68     with np.errstate(all="ignore"):
---> 69         return op(a, b)
    70

```

TypeError: unsupported operand type(s) for /: 'str' and 'int'

During handling of the above exception, another exception occurred:

TypeError Traceback (most recent call last)

<ipython-input-72-cc06344e3055> in <module>

```

----> 1 rmse = kfold(train,3,0.01,0.01)
      2 rmse

```

<ipython-input-71-75fa16f74136> in kfold(xtrain, k, alpha, lamda)

```

    27 def kfold(xtrain,k,alpha,lamda):
    28     errors = []
---> 29     xtrain_k = divide(xtrain,k)
    30     kdata = k_traintest(xtrain_k,k)
    31     for i in range(0,k):

```

<ipython-input-71-75fa16f74136> in divide(data, k)

```

    1 def divide(data,k):
----> 2     size_of_k = math.floor(len(data/k))
    3     data_kdivided = []
    4     c = 0
    5     for i in range(0,k):

```

```

~\anaconda3\lib\site-packages\pandas\core\ops\common.py in new_method(self,
↳ other)

```

```

    63         other = item_from_zerodim(other)
    64
---> 65         return method(self, other)
    66
    67     return new_method

```

```

~\anaconda3\lib\site-packages\pandas\core\arraylike.py in __truediv__(self,
↳ other)

```

```

    111     @unpack_zerodim_and_defer("__truediv__")
    112     def __truediv__(self, other):
--> 113         return self._arith_method(other, operator.truediv)

```

```

114
115     @unpack_zerodim_and_defer("__rtruediv__")

~\anaconda3\lib\site-packages\pandas\core\frame.py in _arith_method(self, other,
↳ op)
    5980         self, other = ops.align_method_FRAME(self, other, axis,
↳ flex=True, level=None)
    5981
-> 5982         new_data = self._dispatch_frame_op(other, op, axis=axis)
    5983         return self._construct_result(new_data)
    5984

~\anaconda3\lib\site-packages\pandas\core\frame.py in _dispatch_frame_op(self,
↳ right, func, axis)
    6006         if not is_list_like(right):
    6007             # i.e. scalar, faster than checking np.ndim(right) == 0
-> 6008             bm = self._mgr.apply(array_op, right=right)
    6009             return type(self)(bm)
    6010

~\anaconda3\lib\site-packages\pandas\core\internals\managers.py in apply(self,
↳ f, align_keys, ignore_failures, **kwargs)
    423         try:
    424             if callable(f):
-> 425                 applied = b.apply(f, **kwargs)
    426             else:
    427                 applied = getattr(b, f)(**kwargs)

~\anaconda3\lib\site-packages\pandas\core\internals\blocks.py in apply(self,
↳ func, **kwargs)
    376         """
    377         with np.errstate(all="ignore"):
-> 378             result = func(self.values, **kwargs)
    379
    380         return self._split_op_result(result)

~\anaconda3\lib\site-packages\pandas\core\ops\array_ops.py in
↳ arithmetic_op(left, right, op)
    187         else:
    188             with np.errstate(all="ignore"):
-> 189                 res_values = _na_arithmetic_op(lvalues, rvalues, op)
    190
    191         return res_values

~\anaconda3\lib\site-packages\pandas\core\ops\array_ops.py in
↳ _na_arithmetic_op(left, right, op, is_cmp)
    147             # will handle complex numbers incorrectly, see GH#32047
    148             raise

```

```

--> 149         result = _masked_arith_op(left, right, op)
      150
      151     if is_cmp and (is_scalar(result) or result is NotImplemented):

~\anaconda3\lib\site-packages\pandas\core\ops\array_ops.py in _
↪ _masked_arith_op(x, y, op)
      109         if mask.any():
      110             with np.errstate(all="ignore"):
--> 111                 result[mask] = op(xrav[mask], y)
      112
      113     result, _ = maybe_upcast_putmask(result, ~mask, np.nan)

```

TypeError: unsupported operand type(s) for /: 'str' and 'int'

[]: